

Elasticsearch 101

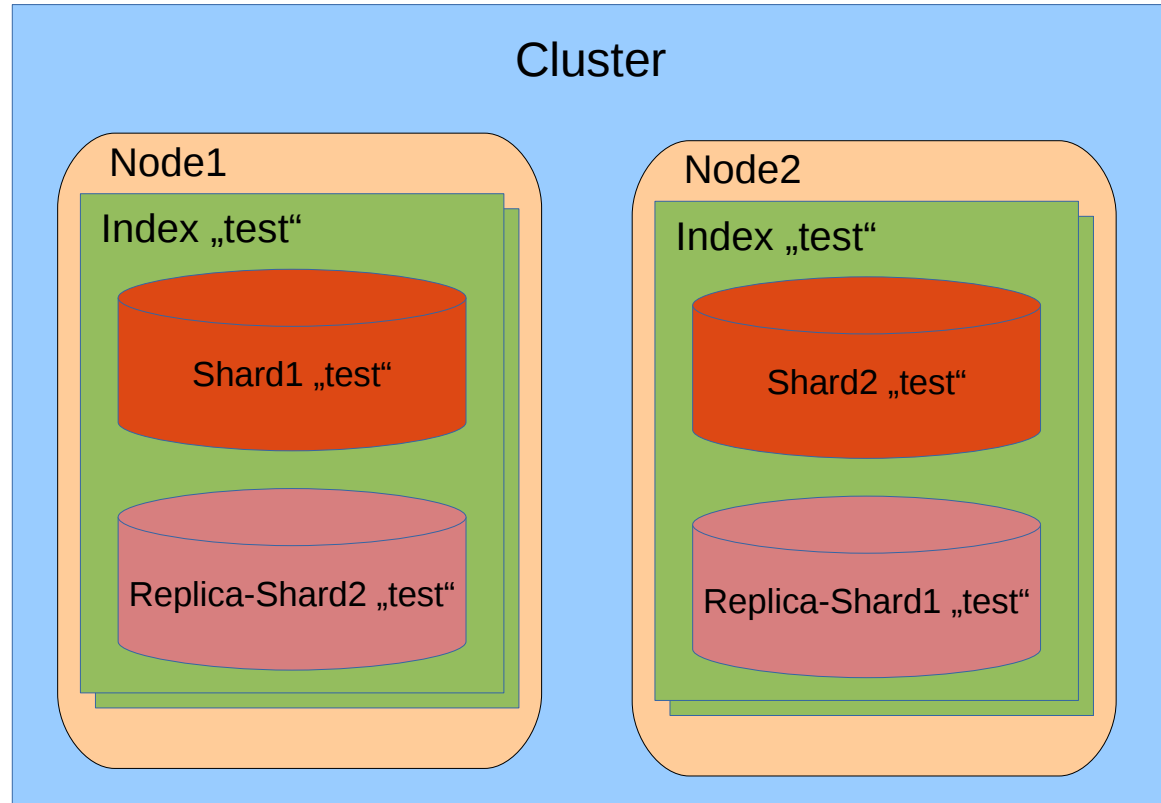


Agenda

- ES Architecture
- CRUD
- Basic Queries
- Similarity Search
- Full-Text Search

Basic Concepts

- Near Realtime
- Cluster
- Node
- Index
- Document
- Shards





Task 1 - CRUD

- Create a new index called „cars“
 - Make sure your cluster is green after this request
 - Set number_of_shards to 2
- Insert a document like this with ID=1:
 - `{"brand":"BMW","series":"5","color":"GREEN","power_kw":100,"marketing_text":"Bringing together BMW's heritage of innovation and design excellence to perform stunningly, the BMW 5 Series grabs attention for all the right reasons."}`
- Find your inserted car by ID



Task 2 – Basic Queries

- Add some more documents to the cars-index, e.g:
 - {"brand":"BMW","series":"3","color":"BLUE","power_kw":150,"marketing_text":"Since 1975, the BMW 3 Series has led the field. With unwavering dedication to the journey, it brings together advanced engineering and expert craftsmanship to create The Ultimate Driving Machine."}
 - {"brand":"VW","series":"Golf","color":"BLACK","power_kw":90,"marketing_text":"The iconic Golf now with pioneering new engines, keener design, new assistance systems and a completely new generation of infotainment systems."}
 - {"brand":"Toyota","series":"Prius","color":"BLACK","power_kw":120,"marketing_text":"Be in your element."}
- Count the number of documents in your index
- Find all BLACK cars
- Find all cars with more than 100kw power

Indices / Mapping / Datatypes

- ES tries to identify the datatype and all fields are indexed by default
- existing field mappings can't be updated!
 - ignore_above can be changed
 - New fields can be added
 - Additional „multi-field“ mappings can be added to existing fields (add additional index to existing field)



Task 3 – Custom Mapping

- Inspect the auto-created cars index mapping
- Delete the cars index
- Create a custom cars index fitting to our data/requirements
- Insert the data from the previous tasks
- Retry the queries from Task 2

Similarity Search

- Elasticsearch provides a **_score** property for each match in a query result.
- In the mapping you can define a **boost** for each datatype-index.
- In a search request you can set a **boost** for a query, which overrides the defined boost in the mapping.
- Similarity Search Query types:
 - More Like This Query
 - Bool Query (should)



Task 4 – Similarity Search

- Build a Bool-Query to find similar cars for a search request:
 - power_kw should be greater than 100, this is most important.
 - color should be BLACK or BLUE, this is mid important.
 - brand should be BMW, this is less important.
- Build a More Like This Query to find similar cars...
 - Search in "marketing_text", "brand" and "series" for “BMW”
 - HINT: set *min_term_freq* and *min_doc_freq* to 1
 - for a existing document in the cars index (_id=1 from Task 1) for “brand”, “color” and “series”
 - HINT: set *min_term_freq* and *min_doc_freq* to 1

How does ES indexing work?

- Es stores data in a inverted index

<https://www.elastic.co/guide/en/elasticsearch/guide/current/inverted-index.html>

- Text is analyzed

<https://www.elastic.co/guide/en/elasticsearch/reference/current/analyzer-anatomy.html>

- 1. Character filters
- 2. Tokenizer
- 3. Token filters



Task 5 – Full-Text Search

- Add an additional type mapping for field “marketing_text” and use the “english” analyzer.
- Search for “innovative car” in “marketing_text” and use the previously created type mapping.
- Also compare the results with the default analyzer used in the default mapping for “marketing_text”.




Bonus Task 1 – Aggregations

- Create a Terms Aggregation for color to get all available colors with the corresponding count of documents.
- Create a Aggregation over “power_kw” to display min/max/average.



Bonus Task 2 – Function Score

- You want to search for `power_kw = 115`, but you want also to score cars which are near by the provided value, but with a lower score.





Bonus Task 3 – add a new Field

- You want to add a new Field “power_ps” which is calculable from power_kw.
 - First update the mapping.
 - Then create a “migrate script” which calculates the new field value from power_kw and re-indexes the documents. HINT: use a painless script.